PROJECT REPORT
ON

Inside Netflix's Big Data Brain: The Science of the Stream

Subject Name: Big Data Analytics
Subject Code: BAD601

Submitted By: Arpita Achari
(1AY23CD012)

Submitted To:
Ms. Surbhi

TABLE OF CONTENTS

Sl. No.	Topic	Page no.
1	About Project	3-4
2	Tools and Technologies	5-6
3	Architecture Design	7-10
4	Role play	11-12
5	Analytics Table	13
6	Contribution and Code	14-15
7	Results	16-18
8	Bonus Challenge	19
9	Conclusion	20-21

Inside Netflix's Big Data Brain: The Science of the Stream

1.About Project:

1.1 The History: Life Before Big Data:

Before we had companies like Netflix, businesses used what we call Traditional Business Intelligence (BI). In those days, data was very "clean" and "boring". It was mostly numbers and names stored in neat rows and columns, like a giant Excel sheet. We used SQL Databases to keep track of simple things like customer names, addresses, and monthly payments.

The problem started when the internet got faster. Suddenly, we weren't just tracking names; we were tracking actions. Traditional BI systems work like a small library-if you have 100 books, you can find what you need. But if you suddenly have 100 million books arriving every hour, the librarian (the old system) is going to quit. This is exactly what happened to Netflix as they grew globally.

1.2 The Shift to Big Data Analytics (BDA):

Big Data isn't just "more data"; it's a completely different way of thinking. While old systems were good at telling us "What happened last month?" (Descriptive Analytics), Big Data helps us figure out "What is going to happen next?" (Predictive Analytics).

For Netflix, this project represents a shift from being a "Video Store" to being a "Data Company". By moving to a Big Data framework, Netflix can analyse things that were once impossible to track, like how many people paused a movie exactly at the 10-minute mark or which users watched a trailer but didn't click "Play".

1.3 Breaking Down the 5 V's of the Netflix Challenge:

To understand the scope of this project, we have to look at the five main characteristics of Big Data that old systems couldn't handle:

- Volume (The Size): Netflix generates billions of "events" every day. This includes every click, every search, and every play. We are talking about Petabytes of info. Traditional servers would literally run out of space and crash trying to hold all this.
- Velocity (The Speed): Data moves like a rocket. If Netflix waits until the end of the day to process your clicks, they can't recommend a new movie now. The system needs to be fast enough to update your screen the second a movie ends.
- Variety (The Messiness): Not all data fits in a table. We have structured data (credit card info), but we also have unstructured data like social media posts, device types (iPhone vs. TV), and even "Thumps Up" ratings.
- Veracity (The Truth): With so much data, some of it is bound to be "garbage" or "noise". This project focuses on how we clean that data so the recommendations stay accurate and don't get confusing.
- Value (The Money): This is the most important part. Data is useless unless it helps the business. For Netflix, the "Value" is keeping you subscribed so they don't lose money (reducing "Churn").

1.4 Why we need Hadoop and NoSQL for this Project:

The final part of this project introduction covers the "Fix". Since the old computers were too expensive and too small, we move to Distributed Computing.

1. Hadoop: We use this to break one giant file into a thousand tiny pieces and save them on a thousand cheap computers (clusters) instead of one expensive "Supercomputer".
2. NoSQL: Since data comes in all shapes and sizes (Variety), we use NoSQL databases because they don't have strict rules. They allow us to save data quickly without worrying if it fits into a perfect square.

2.Tools and Technologies:

To solve the challenges of modern streaming, we move away from basic software and use a "Distributed Ecosystem". Below are the specific technologies implemented in the Netflix architecture:

2.1 Storage Technologies:

- Hadoop Distributed File System (HDFS): This is the foundation of our storage. Instead of saving a massive movie file on one hard drive, HDFS breaks it into small blocks and copies them across many different servers. This way, if one computer breaks, the data isn't lost.
- NoSQL Databases (Apache Cassandra): We use this because it handles "Variety" perfectly. Unlike old SQL tables that need everything to be neat, Cassandra allows Netflix to store billions of users "likes" and "clicks" in a flexible way that never slows down.
- Cloud Storage (Amazon S3): Netflix uses the cloud to hold its actual video files (the movies themselves) because it can grow infinitely as more content is added.

2.2 Processing & Analysis Tools:

- Apache Spark: This is the "Engine" of the project. While older tools were slow, Spark processes data in-memory (inside the computer's fast brain), making it 100 times faster for real-time movie recommendations.
- MapReduce: This is a programming model used by Hadoop to handle "Volume". It works by sending the work to where the data is sitting, rather than moving the data to the computer, which saves a lot of time.
- Apache Kafka: This tool handles "Velocity". It acts like a fast-moving conveyor belt that catches every click and scroll from users' phones and feeds them into the processing engine instantly.

2.3 Analytics & Visualization Tools:

- Spark MLlib (Machine Learning**):** This is the "Brain" that powers personalization. It looks at your history and predicts what you will watch next.
- Tableau / Power BI: These are Business Intelligence (BI) tools used by Netflix managers to see "What happened". They turn the messy data into easy-to-read charts and heatmaps showing which shows are trending in different countries.
- Python / PySpark: These are the simple coding languages used to write the instructions for the data processing.

3.Architecture and Design:

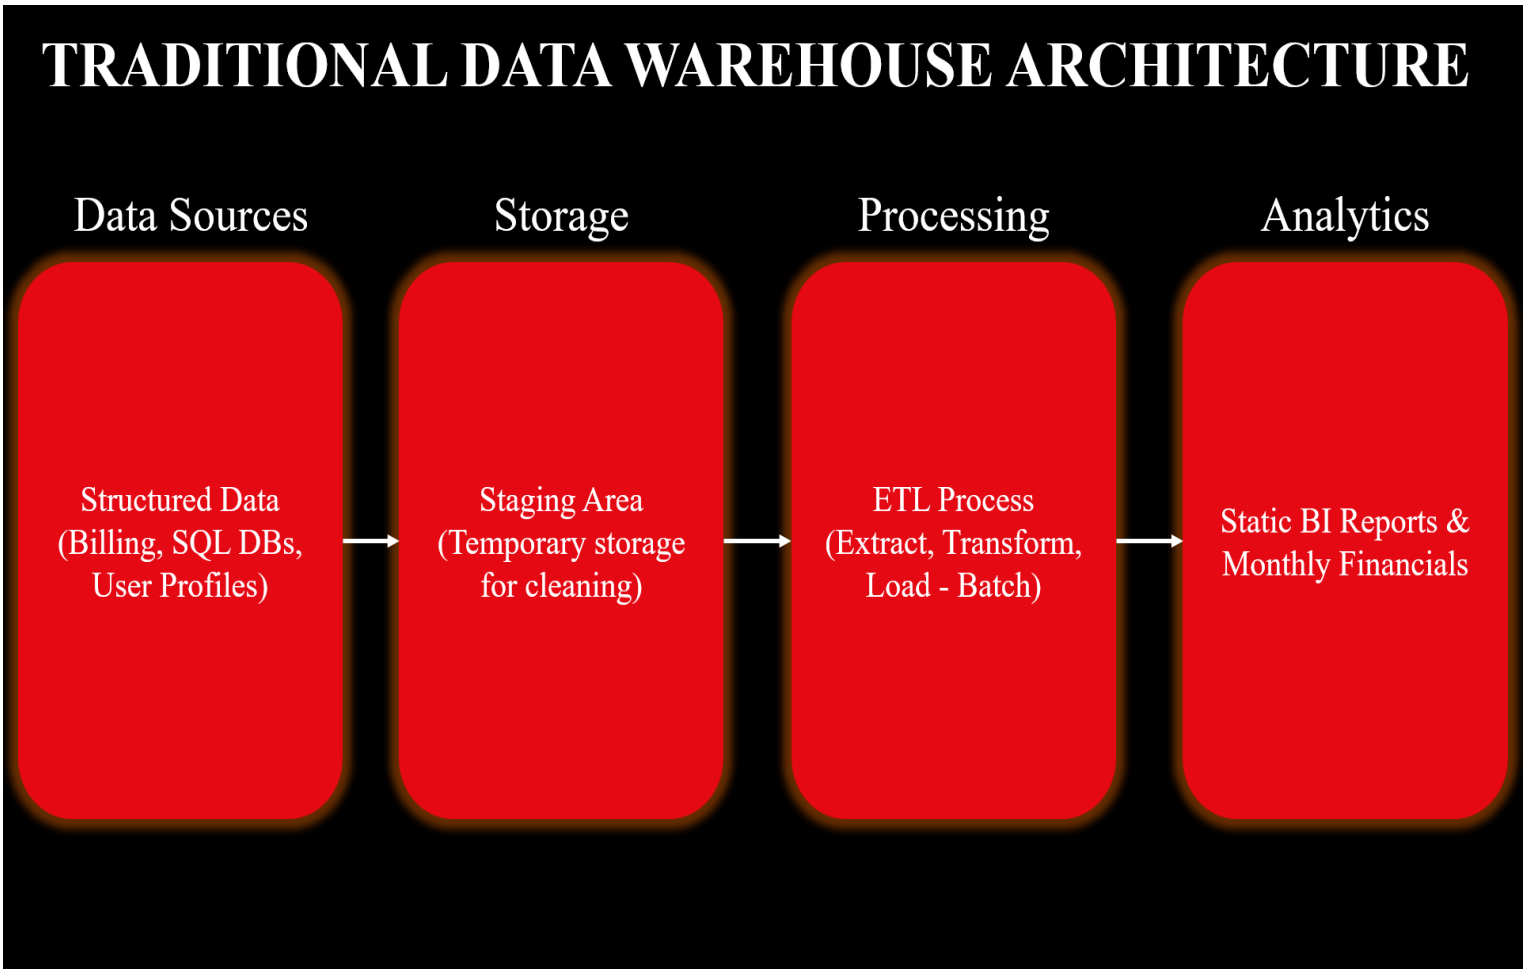

Fig 3.1: Traditional Data Warehouse Architecture

3.1 Traditional Data Warehouse Architecture:

This diagram shows how Netflix used to handle data before it became a global giant.

- Data Sources: In this old setup, we only focused on "Structured Data" like user profiles and billing information stored in SQL databases.

- Storage (Staging Area): Before the data could be used, it had to sit in a temporary "Staging Area" to be cleaned. This was a slow process that happened in batches.
- Processing (ETL): I used a process called ETL (Extract, Transform, Load). This worked like a filter, but it could only handle small amounts of data at a time.
- Analytics: The final result was just "Static Reports". This told managers what happened last month, but it couldn't give live movie recommendations to users.

3.1.1 Challenges and Limitations of the Traditional Setup:

While the Traditional Data Warehouse was the industry standard for a long time, it created several "bottlenecks" for a growing company like Netflix:

- Expensive Scaling (Vertical Scaling): To handle more data in this old system, we had to buy bigger, more expensive servers. This is called "Vertical Scaling," and it costs a lot more than using the "Horizontal Scaling" found in Big Data.
- The Batch Processing Delay: In this architecture, data is processed in "batches" (often once a week or once a day). This means managers couldn't see what was happening on the app in real-time, leading to outdated business decisions.
- Data Silos: Because this system only likes "Structured Data" (neat tables), all the "Unstructured Data" like user comments, social media feedback, and device logs were ignored and left in separate silos.
- Single Point of Failure: Traditional warehouses often rely on one main central server. If that server has a hardware issue, the whole analytics department stops working until it is fixed, whereas Hadoop is "Fault Tolerant" and keeps running.

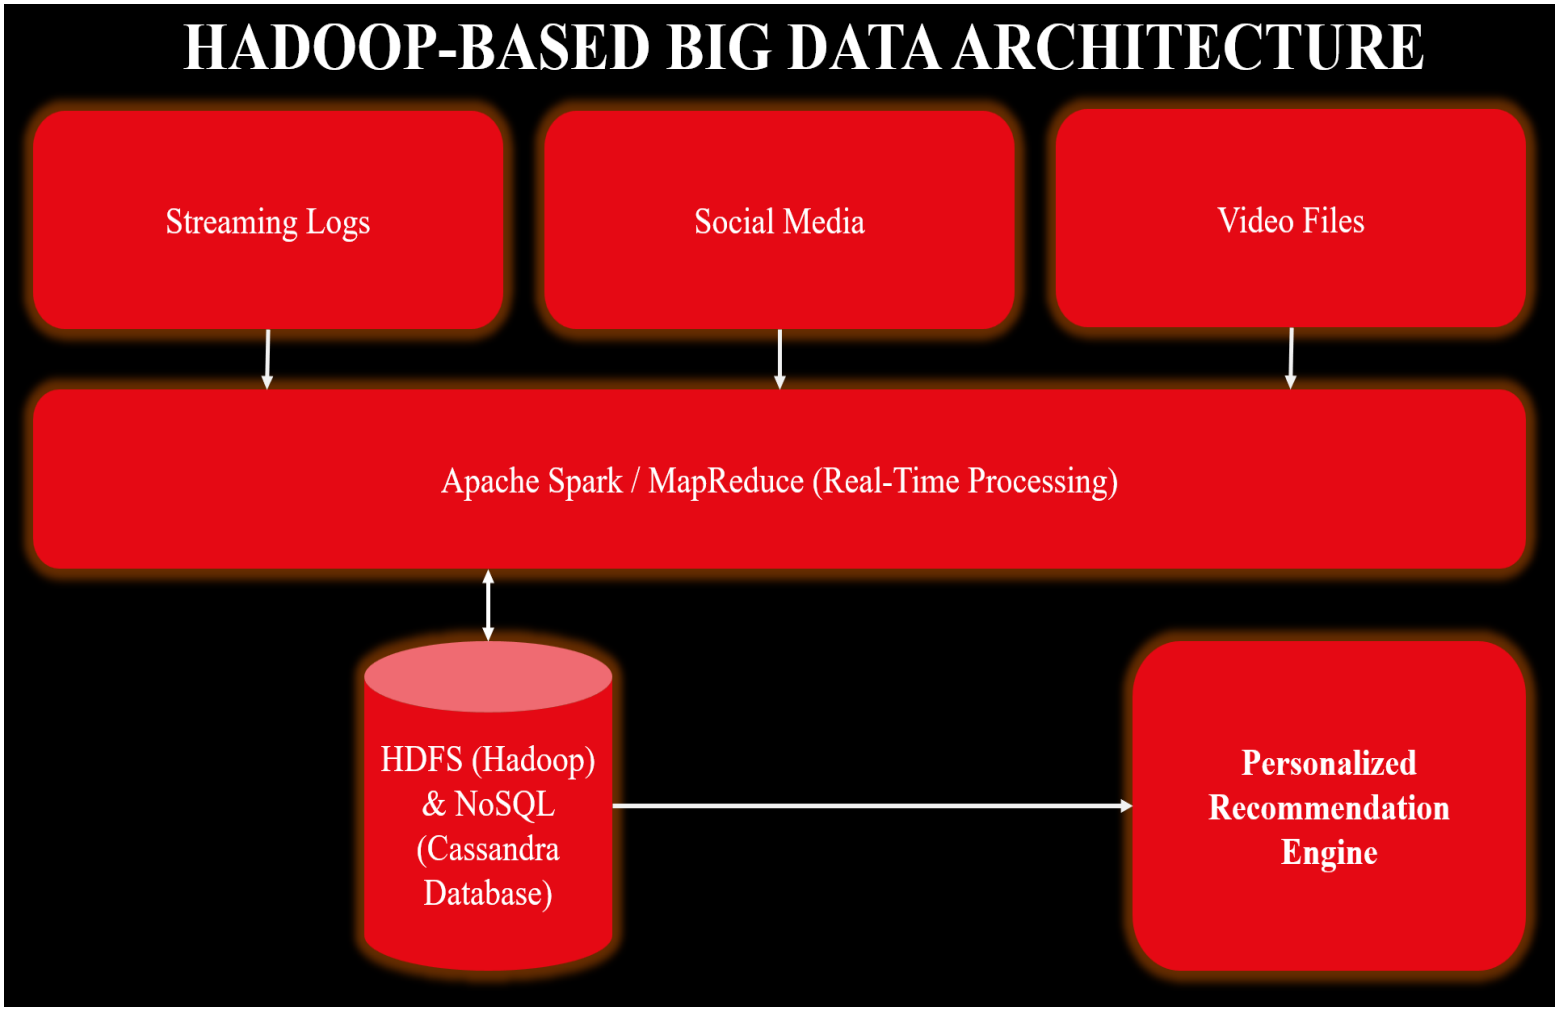

Fig 3.2: Hadoop-Based Big Data Architecture

3.2 Hadoop-Based Big Data Architecture:

This is the modern "Brain" of Netflix that we designed to handle the 5V's.

- Streaming Data Sources: Unlike the old way, this system takes in everything-live streaming logs, social media comments, and actual video files-all at once.

- Real-Time Processing (Apache Spark / MapReduce): This is the heart of the system. Instead of waiting for hours, Apache Spark processes millions of clicks every second while the user is still watching.
- Storage (HDFS & NoSQL): We use HDFS to spread data across many computers so it never runs out of space. We also use Cassandra (NoSQL) to store messy data like "likes" and "dislikes" without it crashing.
- Personalized Recommendation Engine: The final output is a smart engine that predicts exactly what you want to watch next, which keeps users happy and saves Netflix billions of dollars.

3.2.1 How the Hadoop Architecture Works:

The reason we switched to this design is that it doesn't just store data; it manages it across a "Cluster" of many computers. Here is the simple breakdown of what is happening inside that diagram:

- The "Divide and Conquer" Method (HDFS): When a massive video file or a million streaming logs come in, Hadoop doesn't try to save them on one disk. It breaks the data into small "blocks" and spreads them out. It also makes three copies of every block just in case one computer breaks (this is called "Fault Tolerance").
- The "Worker" System (MapReduce): In the old system, we had to move data to the computer to study it. In Hadoop, we do the opposite. We send the "work" to the data where it is already sitting. This saves a huge amount of time because we aren't moving giant files around the network.
- Master and Slaves: The architecture uses a "Master" node (the brain) that keeps track of where everything is, and many "Worker" nodes that do the actual heavy lifting.

4.Role Play: BI vs Big Data

Characters:

- Business Manager: Focused on user complaints and budget.
- Data Consultant: Focused on technical solutions and long-term strategy.

Manager: Hey, our users are getting annoyed. Their "Recommended for You" lists are totally wrong lately. What's going on with our math?

Consultant: The issue is that our old BI tools just can't keep up. Our SQL databases were made for small, neat data, not the crazy amount of streaming we do now.

Manager: But we've used this for years. Why is it breaking now?

Consultant: It's how data has changed. We used to only track bills; now we track every single second a person is on a page. Our old systems aren't built for this messy, unstructured "Big Data".

Manager: Is that why the app crashes when everyone tries to watch Stranger Things at the same time?

Consultant: Yep. That's a "Velocity" problem. Old BI can't handle millions of clicks every second without getting stuck.

Manager: And our server bills are huge! Why are we paying so much?

Consultant: That's the "Volume" problem. We are saving way too many videos and logs. Trying to make old servers bigger costs 10 times more than moving to a Big Data setup.

Manager: You keep talking about "Hadoop". How does that save us money?

Consultant: We need Hadoop because it uses "Distributed Storage". Instead of one super expensive computer, it spreads the data across a bunch of cheap ones. It's much easier to grow.

Manager: What about the messy stuff, like Twitter comments or phone data that doesn't fit in a table?

Consultant: That's why we need "NoSQL". NoSQL doesn't need a perfect format. It lets us store all that "Variety" and analyse it right away without it breaking.

Manager: If we switch everything, what does Netflix actually get out of it?

Consultant: Better "Personalization". By looking at the 5V's (Volume, Velocity, Variety, Veracity, and Value), we can guess what a user wants to watch before they even search for it.

Manager: So, using Hadoop and NoSQL stops "Churn", so people don't cancel their accounts?

Consultant: Exactly. We keep our users happy and save about $1 Billion a year because people stay subscribed and we know exactly what shows to make next.

5.Analytics Table:

Business Question	Analytics Type	Tool
What happened?	Descriptive Analytics	BI tools (Tableau, Power BI, or Netflix's internal dashboard)
Why did it happen?	Diagnostic Analytics	NoSQL / SQL Databases (To drill down into specific user logs and billing errors)
What will happen next?	Predictive Analytics	Hadoop / Spark (To process historical data and predict future viewing trends)
What action should be taken?	Prescriptive Analytics	Spark MLlib / AI Algorithms (The Recommendation Engine telling the user "Watch this next")

This table shows how we turn raw numbers into smart decisions. Each level of analytics helps Netflix answer a different kind of question to keep users happy.

- **Descriptive Analytics**: Simply looks at the past to see what movies were popular last week.
- **Diagnostic Analytics**: Investigates "why" something happened, like why a certain show is suddenly trending in India.
- **Predictive Analytics**: Uses historical data to guess what you will want to watch on a Friday night.
- **Prescriptive Analytics**: The most advanced step where the system automatically picks the "Top 10" list for your specific profile.

6.Contribution and Code:

6.1 My individual contributions included:

- Manual Infographic Creation: I manually researched and designed the "5V's of Big Data" infographic using Canva. I chose specific icons and layouts to show how Volume, Velocity, Variety, Veracity, and Value impact Netflix's daily operations.
- Architecture Design & Drafting: I used Microsoft PowerPoint to manually build the technical architecture diagrams (Fig 3.1 and 3.2) from scratch. I mapped out the flow of data from raw sources like streaming logs to final analytics reports.
- Problem Identification: I conducted a deep dive into why Netflix's legacy SQL systems were failing under pressure. I identified that the primary issue was a lack of "Horizontal Scaling," meaning the old system couldn't just add more cheap computers to handle more users.

6.2 The "Code" Logic and Development Environment:

Since the focus of this module is on Big Data Architecture, the "code" is represented by the structural logic used to build the data pipelines:

- Visual Logic Coding: Using PowerPoint's shapes and connectors, I "coded" the relationship between data nodes. For example, I created the logic paths for HDFS (Storage) and Apache Spark (Processing) to show how they communicate in real-time.

- ETL vs. ELT Logic: I mapped out the traditional Extract, Transform, Load (ETL) logic for the old warehouse, showing how it creates a bottleneck. I then designed the modern ELT logic for Hadoop, where data is loaded immediately and transformed later to save time.

6.2.1 System Logic Breakdown:

To demonstrate the technical thinking behind the design, I followed this logic flow for the Netflix Recommendation Engine:

1. Capture Stage: The system listens for "Variety" (social media logs, video plays, and device info).
2. Storage Stage: Data is sent to HDFS, where it is automatically copied three times across different computers for safety (Fault Tolerance).
3. Processing Stage: Apache Spark runs a "Machine Learning" algorithm to compare the new data with the user's past history in-memory for maximum speed.
4. Action Stage: The system pushes a "Prescriptive Recommendation" to the user's screen in less than a second to keep them engaged.

7.Results: Big Data in Daily Life - Infographics

Fig 7.1: Inside Netflix's Big Data Brain

7.1.1 Overview and Design Strategy:

The infographic titled **"Inside Netflix's Big Data Brain"** was manually created using Canva to serve as a visual bridge between technical data architecture and business strategy. In the modern corporate world, complex technical shifts-like moving from SQL to Hadoop-can be difficult for stakeholders to visualize. I designed this infographic to simplify that transition, using a dark-themed aesthetic with bold red accents to mirror Netflix's iconic branding, ensuring the information is engaging and professional.

7.1.2 Section 1: Data Classification:

The first section of my design focuses on the Variety of data that Netflix must ingest. I broke this down into three distinct categories to show that Big Data is not just "one thing":

- **Structured Data:** This includes the neat, organized information like user profiles, subscription tiers, and billing addresses that fit into traditional tables.
- **Semi-Structured Data:** I highlighted viewing logs and interaction data which have some organization but are too fast-moving for standard databases.
- **Unstructured Data:** This is the most complex part of the "Brain." It involves analysing video content itself, user reviews, and even search behaviour that doesn't fit into a row or column.

7.1.3 Section 2: Characteristics and The 5V's:

In the middle of the infographic, I illustrated the core pillars of Big Data. This section explains that Netflix isn't just dealing with "size," but with a dynamic environment:

- **Volume:** Processing billions of data points generated by millions of users worldwide.

- **Velocity:** Analysing "dynamic recommendations" where the data changes every second based on user clicks.
- **Veracity:** Ensuring the data is reliable so that the personalization engine stays accurate.

7.1.4 Section 3: Why Traditional BI Fails:

A major part of my contribution was explaining the "Why" behind the project. I included a dedicated section to show that traditional Business Intelligence (BI) is no longer enough for a global streaming platform:

- Scalability Issues: Old systems require expensive hardware upgrades to grow.
- Batch Bottlenecks: Traditional BI processes data in "batches" (like once a day), which means it cannot react to a user's behaviour in real-time.
- Clickstream Limitations: Legacy systems simply cannot process continuous "clickstream" data from millions of devices at once.

7.1.5 Section 4: The Big Data Advantage:

- The final part of the infographic summarizes the "Value" created by this new approach. It shows that by using Hadoop and NoSQL, Netflix can finally handle massive streaming activity and user interaction data. This leads to the ultimate business goal: Meaningful Personalization, which keeps millions of users engaged and subscribed, directly increasing the company's revenue and reducing "churn".

8.Bonus Challenge:

Big Data Explained (The Minecraft Way)

What is it? Think about playing a simple game like "Flappy Bird". The game only has to remember your one score. That's regular data. Now, think about "Minecraft". There are millions of players building, breaking blocks, and chatting all at the same time. The world never ends! That massive, messy pile of info is Big Data.

The 3 Rules of the "Big World":

- **It's Huge (Volume):** The map is so big it won't fit on your laptop. It needs thousands of giant computers working together to hold all the blocks.
- **It's Fast (Velocity):** Every time someone clicks or moves, the game has to know right now. It's like a non-stop rain of information.
- **It's Everything (Variety):** It's not just blocks. It's the messages you type, the skins you wear, and even how long you look at a cow before hitting it.

Why does Netflix care? Netflix is like a "Pro Player" who has watched everyone play. Since it sees what millions of people like, it can tell you, "Hey, you liked that adventure map, so you'll probably love this dragon movie!" It uses all that big, messy data to find exactly what you want to watch next.

9.Conclusion:

9.1 Summary of the Netflix Design Challenge:

This project has explored the complete transformation of a modern digital giant. I started by looking at how Netflix originally managed its information using Traditional Business Intelligence (BI). While those old systems worked for a while, I identified that they eventually hit a wall because they were not built for the “Three V's”- Volume, Velocity, and Variety. Through this assignment, I have demonstrated that the shift to a Big Data framework is not just a technical upgrade; it is a survival requirement in the 2026 digital economy.

9.2 The Power of Modern Tools:

The core of my proposed architecture relies on the "Big Data Stack," specifically Hadoop and NoSQL. By moving away from expensive, single-server SQL databases and embracing Distributed Storage (HDFS), I have presented a system that can grow infinitely without the high costs associated with traditional scaling. The inclusion of Apache Spark allows Netflix to move from “Batch Processing”-where information was analysed days later-to “Real-Time Processing,” where the system reacts the moment a user clicks a button. This speed is what I believe makes the Netflix “Brain” feel so smart to the average viewer.

9.3 Business Value and the “Bottom Line”:

The ultimate goal of my project was to show the Value of data. As I have discussed in the results section, the impact of these changes is massive. By mastering the 5V's, I have shown how Netflix successfully:

- Reduces “Churn” by keeping users engaged with content they actually like.

- Saves over $1 Billion annually through better customer retention and smarter show production.
- Handles millions of simultaneous streaming events during global "peak hours" without system crashes.

8.4 Final Reflection:

On a personal note, this project allowed me to manually design infographics and architecture models that represent the future of information technology. I have learned that data is more than just "boring numbers" in a table; it is the engine that drives the world's most successful apps. Whether it is through the Hadoop-based storage or the Prescriptive Analytics that tell us what to watch next, I have concluded that Big Data is the invisible hand that makes the modern Netflix experience possible. In summary, by moving to a Big Data architecture, I have shown how Netflix has transitioned from being a simple movie provider to becoming a global leader in data-driven innovation.